

USED CARS PRICE PREDICTION

Prepared by:

Kateryna Kondratovych

Claudia Słaboń

INTRODUCTION

Out[6]:

ID	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage	Engine	Power	Seats	New_Price	Price
3509	Hyundai i20 Asta Option 1.4 CRDi	Coimbatore	2017	24153	Diesel	Manual	Second	22.54	1396.0	88.73	5.0	NaN	9.18
3332	Mahindra Quanto C6	Mumbai	2013	35000	Diesel	Manual	First	17.21	1493.0	100.00	7.0	NaN	3.49
5383	Tata Indigo LX BSII	Hyderabad	2005	92000	Diesel	Manual	Second	16.10	1405.0	70.00	5.0	NaN	1.10
1891	Maruti Swift Ldi BSIV	Delhi	2014	62000	Diesel	Manual	First	17.80	1248.0	75.00	5.0	NaN	3.75
5757	Toyota Innova Crysta 2.4 VX MT	Kolkata	2017	31000	Diesel	Manual	First	13.68	2393.0	147.80	7.0	21.36 Lakh	16.95

In [7]: X_train.describe()

Out[7]:

	Year	Kilometers_Driven	Mileage	Engine	Power	Seats	Price
count	4815.000000	4.815000e+03	4814.000000	4790.000000	4702.000000	4785.000000	4815.000000
mean	2013.369055	5.915726e+04	18.122254	1619.069102	112.886918	5.274399	9.429570
std	3.286841	1.005633e+05	4.576851	594.356557	52.919245	0.804034	11.246342
min	1998.000000	1.710000e+02	0.000000	624.000000	34.200000	0.000000	0.440000
25%	2011.000000	3.349050e+04	15.170000	1198.000000	76.000000	5.000000	3.500000
50%	2014.000000	5.308000e+04	18.120000	1493.000000	94.000000	5.000000	5.600000
75%	2016.000000	7.300000e+04	21.100000	1989.500000	138.100000	5.000000	9.775000
max	2019.000000	6.500000e+06	33.540000	5461.000000	560.000000	10.000000	160.000000

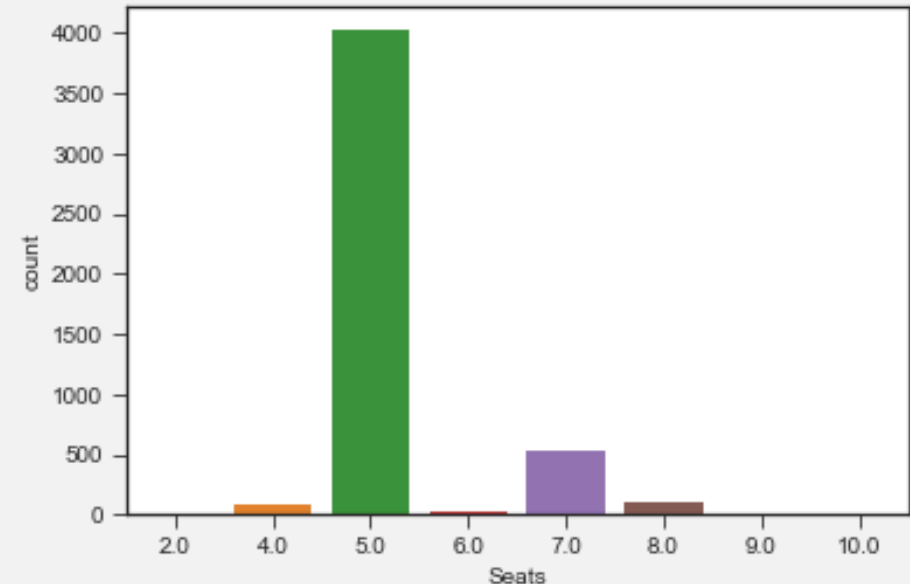
Now it is looking much better, but still we have some 0 values in variables, where we cannot have them (for example 0 Seats) and NA. In next section we will work with this problem

PLAN:
Introduction
Missing Values
Descriptive Analysis
Variable Selection and Transformation
Training/test data division, cross-validation
Usage of Chosen Method on test data
Summary and Conclusions

MISSING VALUES

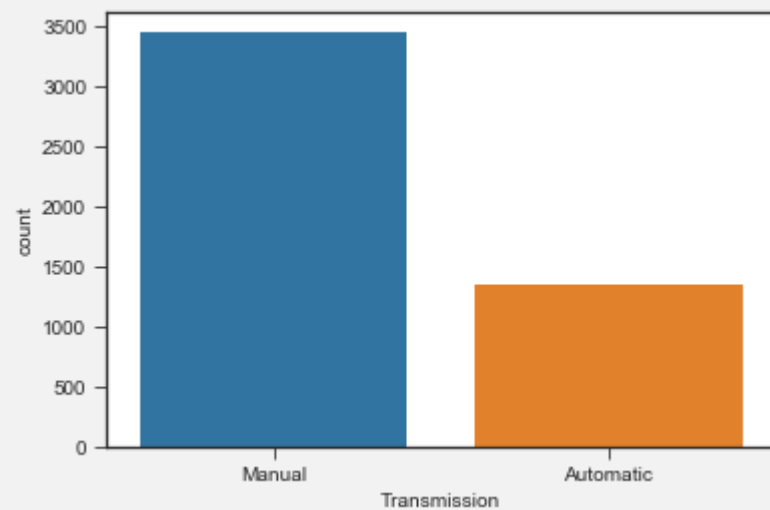
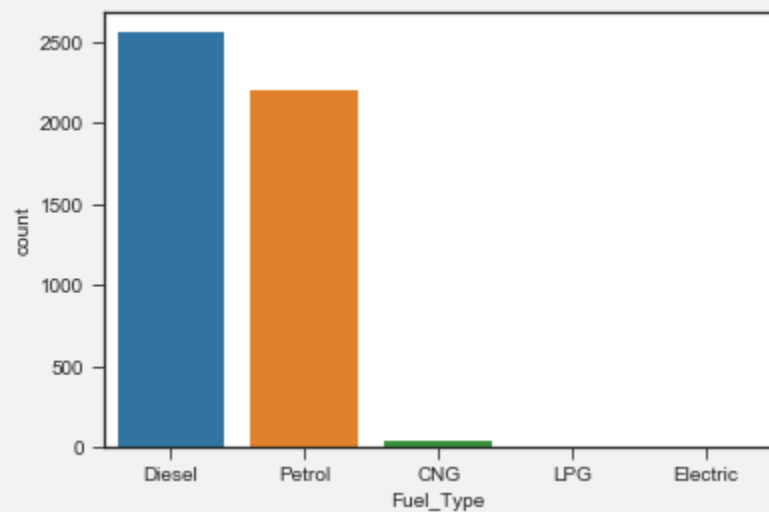
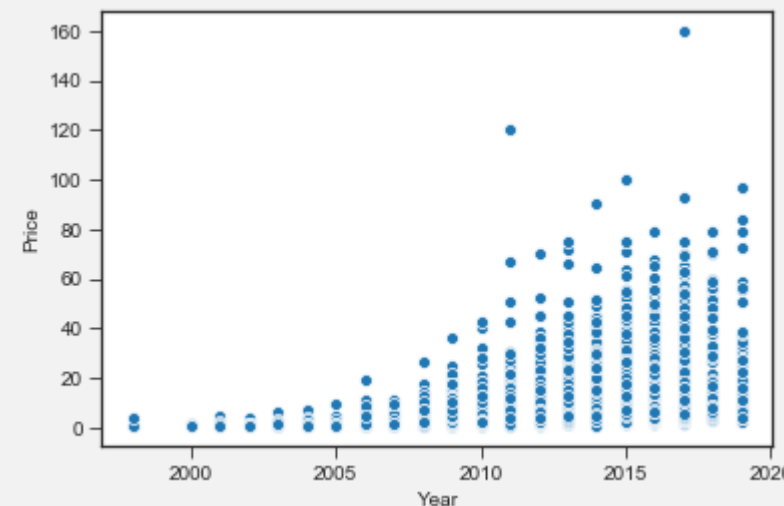
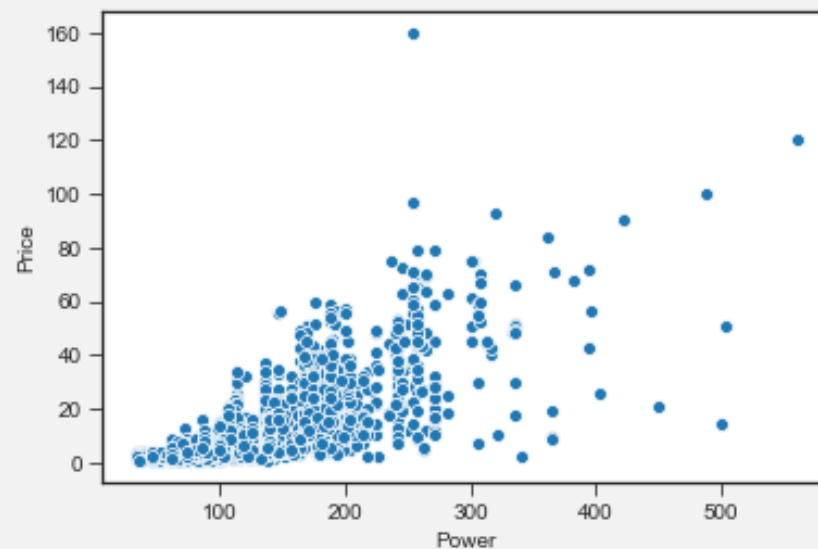
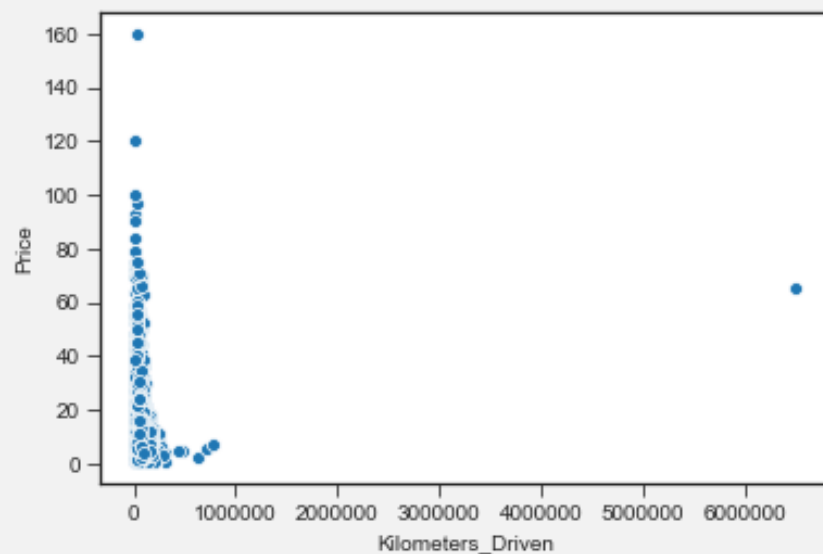
```
Name          0
Location       0
Year           0
Kilometers_Driven 0
Fuel_Type      0
Transmission   0
Owner_Type     0
Mileage        1
Engine         25
Power          113
Seats          30
New_Price     4155
Price          0
dtype: int64
```

```
Percentage of NA's:
Name          0.00
Location       0.00
Year           0.00
Kilometers_Driven 0.00
Fuel_Type      0.00
Transmission   0.00
Owner_Type     0.00
Mileage        0.02
Engine         0.52
Power          2.35
Seats          0.62
New_Price     86.29
Price          0.00
dtype: float64
```

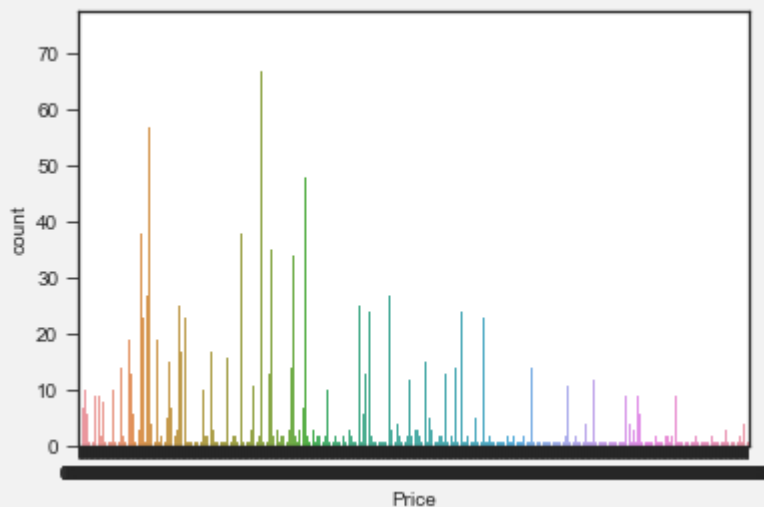


```
by_brand3 = X_train.groupby(['Brand','Brand2','Brand3'])
by_brand2 = X_train.groupby(['Brand','Brand2'])
by_brand1 = X_train.groupby(['Brand'])
def impute_mean(series):
    return series.fillna(series.mean())
X_train.Mileage = by_brand3.Mileage.transform(impute_mean)
X_train[X_train['Mileage'].isnull()]
```

DESCRIPTIVE ANALYSIS

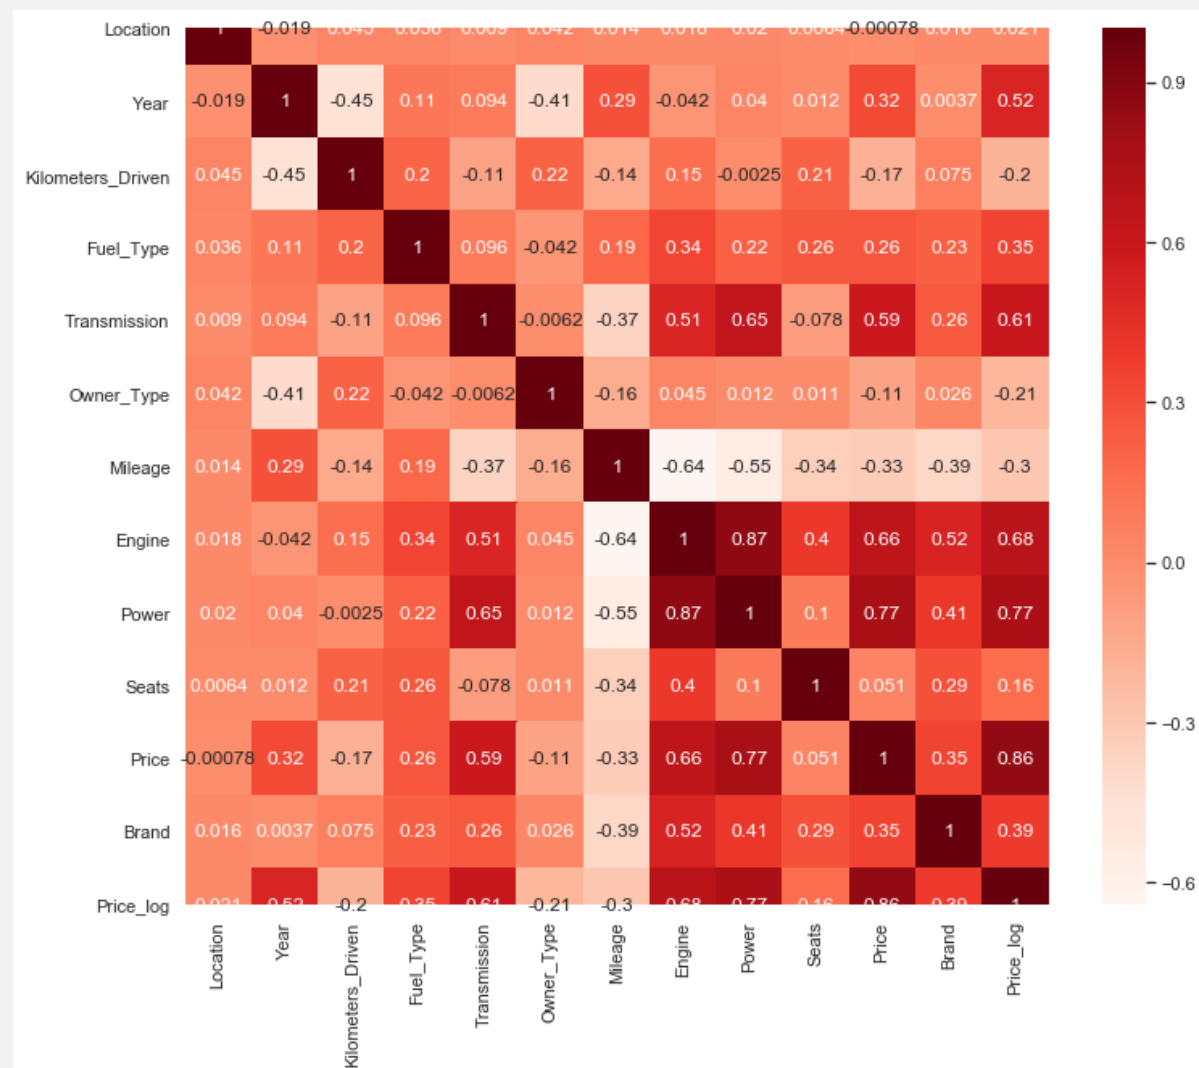


VARIABLE SELECTION AND TRANSFORMATION



Skewness: 3.0255877459369818

Year 0.515176
 Transmission 0.613594
 Engine 0.684915
 Power 0.766890
 Price 0.856055
 Brand 0.387337
 Price_log 1.000000



LINEAR REGRESSION

```
In [57]: import random
from sklearn.model_selection import KFold
for z in range(10):
    trainRes = []
    valRes = []
    kf = KFold(n_splits=5, shuffle=True, random_state=random.randint(0,10000))
    for train, test in kf.split(X_train2.index.values):
        mod = sm.GLM.from_formula(formula="Price_log ~ Transmission + Power +Year", data=X_train_f.iloc[train], family=sm.famil
        res = mod.fit()
        predsTrain = res.predict()
        preds = res.predict(X_train_f.iloc[test])
        trainRes.append(metrics.mean_squared_error(X_train_f.iloc[train].Price_log, predsTrain))
        valRes.append(metrics.mean_squared_error(X_train_f.iloc[test].Price_log, preds))
print("Train AUC:", np.mean(trainRes), "Valid MSE:", np.mean(valRes))
```

```
Train AUC: 0.1275945720645234 Valid MSE: 0.12831727597437734
Train AUC: 0.12761032812077613 Valid MSE: 0.12818194157573876
Train AUC: 0.1276075343983899 Valid MSE: 0.12821580511469835
Train AUC: 0.12756324668486982 Valid MSE: 0.1285784132328012
Train AUC: 0.12762263267644886 Valid MSE: 0.12806959913901234
Train AUC: 0.12757056249161258 Valid MSE: 0.1285375786235907
Train AUC: 0.12757353786277217 Valid MSE: 0.1285156784452455
Train AUC: 0.1276057067005062 Valid MSE: 0.12821169445385755
Train AUC: 0.1275354483536078 Valid MSE: 0.12885624213058575
Train AUC: 0.12761289075293686 Valid MSE: 0.12815004727379206
```

```
import random
from sklearn.model_selection import KFold
for z in range(10):
    trainRes = []
    valRes = []
    kf = KFold(n_splits=5, shuffle=True, random_state=random.randint(0,10000))
    for train, test in kf.split(X_train2.index.values):
        mod = sm.GLM.from_formula(formula="Price_log ~ Transmission + Power +Year+Fuel_Type+Brand", data=X_train_f.iloc[train],
        res = mod.fit()
        predsTrain = res.predict()
        preds = res.predict(X_train_f.iloc[test])
        trainRes.append(metrics.mean_squared_error(X_train_f.iloc[train].Price_log, predsTrain))
        valRes.append(metrics.mean_squared_error(X_train_f.iloc[test].Price_log, preds))
print("Train AUC:", np.mean(trainRes), "Valid MSE:", np.mean(valRes))
```

```
Train AUC: 0.07302763553890937 Valid MSE: 0.1011634980011999
Train AUC: 0.07298551321078126 Valid MSE: 0.10160697994625134
Train AUC: 0.07294813595158618 Valid MSE: 0.10217944142961957
Train AUC: 0.0730209122767681 Valid MSE: 0.10208639132572044
Train AUC: 0.07295204710083966 Valid MSE: 0.10187526789246867
Train AUC: 0.07300728947394422 Valid MSE: 0.10162813289375865
Train AUC: 0.07297449342117264 Valid MSE: 0.1016724216548834
Train AUC: 0.07293803779894117 Valid MSE: 0.10204718606630803
Train AUC: 0.07294829575979596 Valid MSE: 0.10082738300318998
Train AUC: 0.07296784463098271 Valid MSE: 0.10241988786548735
```

KNN

```
kf = KFold(n_splits=5, shuffle=True, random_state=2018)
probs = []
indicies = []
mse = []
r=[]

n_neighbors = 25
clf = neighbors.KNeighborsRegressor(n_neighbors, n_jobs=-1, p=2) #p=2 means Euclidean metrics
for train, test in kf.split(X_train2.index.values):
    clf.fit(X_train2.iloc[train][features].values, X_train_f.iloc[train]["Price_log"].values)
    prob = clf.predict(X_train2.iloc[test][features].values)
    probs.append(prob)
    indicies.append(test)
    mse.append(metrics.mean_squared_error(X_train_f.iloc[test]["Price_log"].values, prob))
    r.append(metrics.r2_score(X_train_f.iloc[test]["Price_log"].values, prob))
print("MSE:")
print(np.mean(mse))
print(mse)
print("\nR_squared:")
print(np.mean(r))
print(r)
```

MSE:
0.1007171235197765
[0.09273168946386275, 0.10477214524931174, 0.09160631982743386, 0.10476296344670001, 0.10971249961157416]

R_squared:
0.8664663862072419
[0.8786857079752665, 0.8622284346216833, 0.8825862757616265, 0.8664687179029944, 0.8423627947746393]

Also, we tried n_neighbors = 50

```
kf = KFold(n_splits=5, shuffle=True, random_state=2018)
probs = []
indicies = []
mse = []
r=[]

n_neighbors = 50
clf = neighbors.KNeighborsRegressor(n_neighbors, n_jobs=-1, p=2)
for train, test in kf.split(X_train2.index.values):
    clf.fit(X_train2.iloc[train][features].values, X_train_f.iloc[train]["Price_log"].values)
    prob = clf.predict(X_train2.iloc[test][features].values)
    probs.append(prob)
    indicies.append(test)
    mse.append(metrics.mean_squared_error(X_train_f.iloc[test]["Price_log"].values, prob))
    r.append(metrics.r2_score(X_train_f.iloc[test]["Price_log"].values, prob))
print("MSE:")
print(np.mean(mse))
print(mse)
print("\nR_squared:")
print(np.mean(r))
print(r)
```

MSE:
0.1069828207847455
[0.09982755615754128, 0.10968818439624384, 0.09690989012821222, 0.113210916690616, 0.11527755655111413]

R_squared:
0.8582046037599147
[0.8694026888776691, 0.8557640217080957, 0.8757885794678831, 0.8557009237260992, 0.8343668050198262]

```
kf = KFold(n_splits=5, shuffle=True, random_state=2018)
probs = []
indicies = []
mse = []
r=[]

n_neighbors = 25
clf = neighbors.KNeighborsRegressor(n_neighbors, n_jobs=-1, p=1) #p=1 means Manhattan metrics
for train, test in kf.split(X_train2.index.values):
    clf.fit(X_train2.iloc[train][features].values, X_train_f.iloc[train]["Price_log"].values)
    prob = clf.predict(X_train2.iloc[test][features].values)
    probs.append(prob)
    indicies.append(test)
    mse.append(metrics.mean_squared_error(X_train_f.iloc[test]["Price_log"].values, prob))
    r.append(metrics.r2_score(X_train_f.iloc[test]["Price_log"].values, prob))
print("MSE:")
print(np.mean(mse))
print(mse)
print("\nR_squared:")
print(np.mean(r))
print(r)
```

MSE:
0.1004278068984125
[0.09387832418601454, 0.10391610765212066, 0.09025277473552533, 0.1048334485134056, 0.10925837940499644]

R_squared:
0.8668510088432976
[0.8771856470971237, 0.8633540929682512, 0.8843211426405158, 0.8663788773618696, 0.8430152841487281]

SVR

LINEAR

```
from sklearn.svm import SVR
#data have been previously normalized
features = X_train2.columns.tolist()
X_train2.iloc[test][features]
kf = KFold(n_splits=5)
mses = []
r=[]
clf = SVR(C=1, cache_size=500, kernel='linear',
          max_iter=-1, tol=0.001, verbose=False)

for train, test in kf.split(X_train2.index.values):

    clf.fit(X_train2.iloc[train][features].values, X_train_f.iloc[train]["Price_log"].values)
    prob = clf.predict(X_train2.iloc[test][features].values)
    mses.append(metrics.mean_squared_error(X_train_f.iloc[test]["Price_log"].values, prob))
    r.append(metrics.r2_score(X_train_f.iloc[test]["Price_log"].values, prob))

print("MSE:")
print(np.mean(mses))
print(mses)
print("\nR_squared:")
print(np.mean(r))
print(r)
```

MSE:
0.12948012299206738
[0.15874197860758177, 0.11977196275956149, 0.11974562704065317, 0.12674426973800326, 0.12239677681453726]

R_squared:
0.8284736793502641
[0.7797701635591572, 0.838381445566693, 0.8500766896431577, 0.8339830355795458, 0.8401570624027663]

MSE:
0.10451219617870751
[0.11117407194649125, 0.09760624066298232, 0.09637780353971089, 0.11502580685946663, 0.1023770578848865]

R_squared:
0.8618045064796229
[0.8457632448832985, 0.8682915504082667, 0.8793335530600264, 0.8493325550393995, 0.8663016290071237]

RBF 2

POLYNOMIAL

```
features = X_train2.columns.tolist()
X_train2.iloc[test][features]
kf = KFold(n_splits=5)
mses = []
r=[]
clf = SVR(C=1, cache_size=500, degree=2, kernel='poly',
          max_iter=-1, tol=0.001, verbose=False)

for train, test in kf.split(X_train2.index.values):

    clf.fit(X_train2.iloc[train][features].values, X_train_f.iloc[train]["Price_log"].values)
    prob = clf.predict(X_train2.iloc[test][features].values)
    mses.append(metrics.mean_squared_error(X_train_f.iloc[test]["Price_log"].values, prob))
    r.append(metrics.r2_score(X_train_f.iloc[test]["Price_log"].values, prob))

print("MSE:")
print(np.mean(mses))
print(mses)
print("\nR_squared:")
print(np.mean(r))
print(r)
```

MSE:
0.10769943339359962
[0.11920064214864541, 0.09910466250460317, 0.099796650983007, 0.11543015550171343, 0.1049650558300291]

R_squared:
0.8575350210376651
[0.8346276255700784, 0.8662696016450139, 0.8750531050890122, 0.8488029158352595, 0.8629218570489615]

RBF

```
features = X_train2.columns.tolist()
X_train2.iloc[test][features]
kf = KFold(n_splits=5)
mses = []
r=[]
clf = SVR(C=1, cache_size=500, degree=3, kernel='rbf',
          max_iter=-1, tol=0.001, verbose=False)

for train, test in kf.split(X_train2.index.values):

    clf.fit(X_train2.iloc[train][features].values, X_train_f.iloc[train]["Price_log"].values)
    prob = clf.predict(X_train2.iloc[test][features].values)
    mses.append(metrics.mean_squared_error(X_train_f.iloc[test]["Price_log"].values, prob))
    r.append(metrics.r2_score(X_train_f.iloc[test]["Price_log"].values, prob))

print("MSE:")
print(np.mean(mses))
print(mses)
print("\nR_squared:")
print(np.mean(r))
print(r)
```

MSE:
0.10472341869806204
[0.11038363397901829, 0.09866005222685807, 0.09694393518208066, 0.11520769738031583, 0.10242177472203728]

R_squared:
0.8615383381404491
[0.8468598547770401, 0.866869552323952, 0.8786247478032547, 0.8490943043303879, 0.8662432314676111]

USAGE OF CHOSEN METHOD ON TEST DATA

MSE: 0.0975013790102812

R²: 0.8727525123577311

```
n_neighbors = 25
```

```
clf = neighbors.KNeighborsRegressor(n_neighbors, n_jobs=-1, p=1)
```

```
clf.fit(X_test2, Y)
```

```
y_pred=clf.predict(X_test2)
```

```
print(metrics.mean_squared_error(Y, y_pred))
```

```
print(metrics.r2_score(Y, y_pred))
```

Dependent variable: Price_log - logarithm of Price Variable

Independent variables: Transmission, Year, Power

By R² and MSE we chose kNN with n_neighbors = 25 and Manhattan metric

CONCLUSIONS

- Best Model: kNN with `n_neighbors = 25` and Manhattan metric
- Validation: MSE: 0.1004278068984125 R_squared: 0.8668510088432976
- Final: MSE: 0.0975013790102812 R_squared: 0.8727525123577311